

Feature-Value Scoring

Field ownership & usage reference — who scores what (human vs. AI), for Features and Defects

Source of truth: **feature-value-scoring.xlsx** Data Dictionary sheet + the **feature-value-scoring** skill role contracts. If this document and the Data Dictionary ever disagree, the Data Dictionary wins.

Human — judgment / triage	AI (Research agent) — evidence-scored	Shared — human + agent	Calculated — formula, never edit	n/a — not used for this type
----------------------------------	--	-------------------------------	---	-------------------------------------

1 Field ownership matrix

Field	Scale	Features	Defects	What it captures / how it's used
Identity & links				
feature_id	F-### / D-###	Human	Human	Stable join key to the brief (Feature) or bug report (Defect).
feature_name	text	Human	Human	Short human-readable name.
brief_ref	path / URL	Human	Human	Link to the feature brief (F) or the defect/bug report & repro (D).
item_type	Feature Defect	Human	Human	Selects the scoring path. Blank legacy rows default to Feature.
Value dimensions (1–10, scored from evidence)				
reach	1–10	AI / Human	AI / Human	F: share of users affected × frequency. D: share of customers who hit the bug (feeds the defect reach factor).
user_value	1–10	AI / Human	n/a	Depth of pain relieved per affected user. Not used by the defect formula.
business_value	1–10	AI / Human	n/a	Revenue / retention / cost / market-access impact. Defect: n/a.
time_sensitivity	1–10	AI / Human	n/a	Decay of value if delayed. Defect: n/a.
confidence	1–10	AI / Human	AI / Human	F: strength of evidence behind value scores. D: strength of the repro/evidence. Same multiplier map for both.
Human-only judgment				
strategic_alignment	1–10	Human only	n/a	Fit with product vision — deliberately human-owned. Not in the defect formula.
effort	1–10	Human (eng)	Human — optional	Total delivery cost. Divides the Feature score. Collectible for defect fix-planning but NOT in the defect formula.
risk	1–10	Human only	Human — optional	Technical / market / reversibility risk. In the Feature formula; not in the defect formula.
severity	1–5	n/a	Human (triage) only	Defect only. Drives the defect priority directly. NEVER agent-scored — a triage judgment call. Stays 1–5 (industry Sev1–Sev4).
Provenance & workflow				

Field	Scale	Features	Defects	What it captures / how it's used
evidence_sources	text	AI / Human	AI / Human	Citations for every score. Mandatory for agent-scored rows; no citation → confidence capped at 4.
rationale_notes	text	AI / Human	AI / Human	Reasoning behind judgment calls and any dissent.
scored_by	Human Market Research Agent Hybrid	AI / Human	AI / Human	Provenance of the scores. 'Hybrid' when an agent revises human scores.
date_scored	ISO date	AI / Human	AI / Human	When scores were last updated; prioritizer discounts rows > 90 days old.
defect_source	Customer-Reported Internal	n/a (blank)	Human / Support	Defect only. How the defect was found.
status	enum (see §4)	Human + agents	Human + agents	Lifecycle state. Human sets early states; agents advance defined transitions; Cancelled is human-only.
Calculated — formula columns, never written by anyone				
weighted_value	0.00–5.00	formula	blank for Defects	SUMPRODUCT of value dimensions × weights (Feature path only).
confidence_multiplier	0.50–1.00	formula	formula	Lookup from the confidence map.
risk_factor	0.60–1.00	formula	formula	Lookup from the risk map (Feature path).
priority_score	0–100	formula	formula	Final score. Feature vs. defect formula per item_type — see §2.
rank	1–N	formula	formula	Descending rank of priority_score across ALL scored rows — Features and Defects share one list.

2 The two scoring formulas

item_type is the fork: every row is scored by exactly one of these. Both land in the same **priority_score** / **rank** columns, so a critical customer defect and a nice-to-have feature compete in one ranked list — by design.

Feature — weighted value

weighted_value = reach×0.20 + user_value×0.25 + business_value×0.25 + strategic_alignment×0.20 + time_sensitivity×0.10

priority_score = MIN(100, weighted_value × confidence_mult × risk_factor ÷ effort × 20)

Cheap, safe items can outrank strategic bets because effort divides — intentional; the strategic-alignment review flag is the counterweight.

Defect — severity fast-track

priority_score = MIN(100, severity_base × reach_factor × confidence_mult)

Deliberately bypasses weighted_value, strategic_alignment, and the effort divisor — a blocking/data-loss defect shouldn't get cheaper-to-fix bugs discounted below it, nor wait for a critical mass of affected customers. Severity dominates: a low-reach Sev1 still ranks near the top.

3 Config lookup maps

All maps live on the workbook **Config** sheet — read them, never assume. Every dimension is 1–10 except **severity** (1–5).

Confidence → mult	1	2	3	4	5	6	7	8	9	10
→	0.45	0.5	0.58	0.65	0.73	0.8	0.85	0.9	0.95	1
Risk → factor	1	2	3	4	5	6	7	8	9	10
→	1	1	0.95	0.9	0.85	0.8	0.75	0.7	0.65	0.6
Severity → base	1	2	3	4	5					
→	10	30	55	80	100					
Reach → factor (Defect)	1	2	3	4	5	6	7	8	9	10
→	0.75	0.8	0.85	0.9	0.95	1	1.05	1.1	1.15	1.2

4 Who does what — roles & lifecycle

AI — Market Research Agent (RESEARCH role). May write value/evidence fields only. **Features:** reach, user_value, business_value, time_sensitivity, confidence, evidence_sources, rationale_notes, scored_by, date_scored, and status→Scored. **Defects:** reach, confidence, evidence_sources, rationale_notes, scored_by, date_scored, status→Scored. Must cite sources; must not inflate value to offset thin evidence (honest value + low confidence is the correct encoding). **Never writes** strategic_alignment, effort, risk, severity, item_type, defect_source, or any formula column.

AI — Discover & Design roles. Write **no scores at all**. Discovery emits an evidence signal report (and, in scan mode, may create Idea rows with id/name/notes only). Design contributes qualitative reads to a brief. Neither sets a number.

AI — Prioritization agent (PRIORITIZE role). Writes **status only** (Scored→Prioritized | Deferred | Rejected). Flags for a human instead of deciding: severity=5, risk≥9, strategic_alignment≥8 buried at low rank, within ±10% of the funding cut line, or date_scored > 90 days. A 'Won't Fix' call is the human's.

Human. Owns identity (id, name, brief_ref, item_type), the judgment dimensions (strategic_alignment, effort, risk), defect triage (severity, defect_source), and all terminal decisions. Sets initial status and resolves every agent-raised flag. **Cancelled** is human-only, at any time.

Why the split? The actor that frames or estimates an item is never the actor that scores its evidence-based value. Strategic fit, engineering effort, risk appetite, and defect severity require organizational judgment an evidence-gathering agent doesn't hold — so they stay human. Keeping value dimensions evidence-scored and judgment dimensions human-owned is what makes the shared ranked list trustworthy.

Feature status flow	Idea → Brief Drafted → Scored → Prioritized → In Progress → Implemented (or Deferred Rejected)
Defect status flow	Reported → Triaged → Scored → Prioritized → In Progress → Fixed → Verified (or Deferred Won't Fix)

Who advances	Human (early states) · Prioritization agent (Scored→Prioritized) · /implement-spec (Feature build states) · /fix-defect-or-bug (Defect fix states). Cancelled — human only, any point.
---------------------	--